

Pandas

Dr. Noman Islam

Introduction

- It contains data structures and data manipulation tools designed to make data cleaning and analysis fast and easy in Python.
- Pandas is designed for working with tabular or heterogeneous data.
- NumPy, by contrast, is best suited for working with homogeneous numerical array data.

Series

```
In [11]: obj = pd.Series([4, 7, -5, 3])
```

```
In [13]: obj.values
```

```
Out[13]: array([ 4,  7, -5,  3])
```

```
In [14]: obj.index # like range(4)
```

```
Out[14]: RangeIndex(start=0, stop=4, step=1)
```

```
In [15]: obj2 = pd.Series([4, 7, -5, 3], index=['d', 'b', 'a', 'c'])
```

```
In [16]: obj2
```

```
Out[16]:
```

```
d      4
```

```
b      7
```

```
a     -5
```

```
c      3
```

```
dtype: int64
```

```
In [17]: obj2.index
```

```
Out[17]: Index(['d', 'b', 'a', 'c'], dtype='object')
```

Using indexes to access data

```
In [18]: obj2['a']
```

```
Out[18]: -5
```

```
In [19]: obj2['d'] = 6
```

```
In [20]: obj2[['c', 'a', 'd']]
```

```
Out[20]:
```

```
c      3
```

```
a     -5
```

```
d      6
```

```
dtype: int64
```

Filtering

```
In [21]: obj2[obj2 > 0]
```

```
Out[21]:
```

```
d      6
```

```
b      7
```

```
c      3
```

```
dtype: int64
```

```
In [22]: obj2 * 2
```

```
Out[22]:
```

```
d     12
```

```
b     14
```

```
a    -10
```

```
c      6
```

```
dtype: int64
```

```
In [24]: 'b' in obj2
```

```
Out[24]: True
```

Creating series from dictionary

```
In [26]: sdata = {'Ohio': 35000, 'Texas': 71000, 'Oregon': 16000, 'Utah': 5000}
```

```
In [27]: obj3 = pd.Series(sdata)
```

```
In [28]: obj3
```

```
Out[28]:
```

```
Ohio      35000
Oregon     16000
Texas      71000
Utah       5000
dtype: int64
```

```
In [29]: states = ['California', 'Ohio', 'Oregon', 'Texas']
```

```
In [30]: obj4 = pd.Series(sdata, index=states)
```

```
In [31]: obj4
```

```
Out[31]:
```

```
California    NaN
Ohio          35000.0
Oregon        16000.0
Texas         71000.0
dtype: float64
```

Checking for null values

```
In [32]: pd.isnull(obj4)
```

```
Out[32]:
```

```
California    True
```

```
Ohio          False
```

```
Oregon        False
```

```
Texas         False
```

```
dtype: bool
```

```
In [33]: pd.notnull(obj4)
```

```
Out[33]:
```

```
California    False
```

```
Ohio          True
```

```
Oregon        True
```

```
Texas         True
```

```
dtype: bool
```

Adding two series

```
In [35]: obj3  
Out[35]:  
Ohio      35000  
Oregon     16000  
Texas      71000  
Utah       5000  
dtype: int64
```

```
In [36]: obj4  
Out[36]:  
California      NaN  
Ohio            35000.0  
Oregon          16000.0  
Texas           71000.0  
dtype: float64
```

```
In [37]: obj3 + obj4  
Out[37]:  
California      NaN  
Ohio            70000.0  
Oregon          32000.0  
Texas          142000.0  
Utah            NaN
```

Dataframe

```
data = {'state': ['Ohio', 'Ohio', 'Ohio', 'Nevada', 'Nevada', 'Nevada'],  
        'year': [2000, 2001, 2002, 2001, 2002, 2003],  
        'pop': [1.5, 1.7, 3.6, 2.4, 2.9, 3.2]}  
frame = pd.DataFrame(data)
```

```
In [51]: frame2['state']
```

```
Out[51]:
```

```
one      Ohio  
two      Ohio  
three    Ohio  
four     Nevada  
five     Nevada  
six      Nevada
```

```
Name: state, dtype: object
```

```
In [46]: frame.head()
```

```
Out[46]:
```

	pop	state	year
0	1.5	Ohio	2000
1	1.7	Ohio	2001
2	3.6	Ohio	2002
3	2.4	Nevada	2001
4	2.9	Nevada	2002

```
In [48]: frame2 = pd.DataFrame(data, columns=['year', 'state', 'pop', 'debt'],
.....:                          index=['one', 'two', 'three', 'four',
.....:                                'five', 'six'])
```

```
In [51]: frame2['state']
```

```
Out[51]:
```

one	Ohio
two	Ohio
three	Ohio
four	Nevada
five	Nevada
six	Nevada

```
Name: state, dtype: object
```

```
In [52]: frame2.year
```

```
Out[52]:
```

one	2000
two	2001
three	2002
four	2001
five	2002
six	2003

```
In [53]: frame2.loc['three']
```

```
Out[53]:
```

year	2002
state	Ohio
pop	3.6
debt	NaN

```
Name: three, dtype: object
```

Modifying column values

```
In [54]: frame2['debt'] = 16.5
```

```
In [55]: frame2
```

```
Out[55]:
```

	year	state	pop	debt
one	2000	Ohio	1.5	16.5
two	2001	Ohio	1.7	16.5
three	2002	Ohio	3.6	16.5
four	2001	Nevada	2.4	16.5
five	2002	Nevada	2.9	16.5
six	2003	Nevada	3.2	16.5

```
In [56]: frame2['debt'] = np.arange(6.)
```

Adding / deleting a new column

```
In [61]: frame2['eastern'] = frame2.state == 'Ohio'
```

```
In [62]: frame2
```

```
Out[62]:
```

	year	state	pop	debt	eastern
one	2000	Ohio	1.5	NaN	True
two	2001	Ohio	1.7	-1.2	True
three	2002	Ohio	3.6	NaN	True
four	2001	Nevada	2.4	-1.5	False
five	2002	Nevada	2.9	-1.7	False
six	2003	Nevada	3.2	NaN	False

```
In [63]: del frame2['eastern']
```

```
In [64]: frame2.columns
```

```
Out[64]: Index(['year', 'state', 'pop', 'debt'], dtype='object')
```

Transpose a dataframe

```
In [68]: frame3.T
```

```
Out[68]:
```

	2000	2001	2002
Nevada	NaN	2.4	2.9
Ohio	1.5	1.7	3.6

Slicing in dataframe

```
In [70]: pdata = {'Ohio': frame3['Ohio'][:-1],  
.....:          'Nevada': frame3['Nevada'][:2]}
```

```
In [71]: pd.DataFrame(pdata)
```

```
Out[71]:
```

	Nevada	Ohio
2000	NaN	1.5
2001	2.4	1.7

Values attribute

As with Series, the `values` attribute returns the data contained in the DataFrame as a two-dimensional ndarray:

```
In [74]: frame3.values
```


Table 5-1. Possible data inputs to DataFrame constructor

Type	Notes
2D ndarray	A matrix of data, passing optional row and column labels
dict of arrays, lists, or tuples	Each sequence becomes a column in the DataFrame; all sequences must be the same length
NumPy structured/record array	Treated as the “dict of arrays” case
dict of Series	Each value becomes a column; indexes from each Series are unioned together to form the result’s row index if no explicit index is passed
dict of dicts	Each inner dict becomes a column; keys are unioned to form the row index as in the “dict of Series” case
List of dicts or Series	Each item becomes a row in the DataFrame; union of dict keys or Series indexes become the DataFrame’s column labels
List of lists or tuples	Treated as the “2D ndarray” case
Another DataFrame	The DataFrame’s indexes are used unless different ones are passed
NumPy MaskedArray	Like the “2D ndarray” case except masked values become NA/missing in the DataFrame result

Indexing, selection and filtering

```
In [117]: obj = pd.Series(np.arange(4.), index=['a', 'b', 'c', 'd'])
```

```
In [121]: obj[2:4]
```

```
Out[121]:
```

```
c    2.0
```

```
d    3.0
```

```
dtype: float64
```

```
In [122]: obj[['b', 'a', 'd']]
```

```
Out[122]:
```

```
b    1.0
```

```
a    0.0
```

```
d    3.0
```

```
dtype: float64
```

```
In [124]: obj[obj < 2]
```

```
Out[124]:
```

```
a    0.0
```

```
b    1.0
```

```
dtype: float64
```

Dropping entries from axis

```
In [105]: obj = pd.Series(np.arange(5.), index=['a', 'b', 'c', 'd', 'e'])
```

```
In [106]: obj
```

```
Out[106]:
```

```
a    0.0
```

```
b    1.0
```

```
c    2.0
```

```
d    3.0
```

```
e    4.0
```

```
dtype: float64
```

```
In [107]: new_obj = obj.drop('c')
```

```
In [108]: new_obj
```

```
Out[108]:
```

```
a    0.0
```

```
b    1.0
```

```
d    3.0
```

```
e    4.0
```

```
dtype: float64
```

Axis parameter

```
In [113]: data.drop('two', axis=1)
```

```
Out[113]:
```

	one	three	four
Ohio	0	2	3
Colorado	4	6	7
Utah	8	10	11
New York	12	14	15

```
In [114]: data.drop(['two', 'four'], axis='columns')
```

```
Out[114]:
```

	one	three
Ohio	0	2
Colorado	4	6
Utah	8	10
New York	12	14

Indexing and slicing

```
In [119]: obj['b']  
Out[119]: 1.0
```

```
In [120]: obj[1]  
Out[120]: 1.0
```

```
In [121]: obj[2:4]  
Out[121]:  
c      2.0  
d      3.0  
dtype: float64
```

```
In [122]: obj[['b', 'a', 'd']]  
Out[122]:  
b      1.0  
a      0.0  
d      3.0  
dtype: float64
```

```
In [123]: obj[[1, 3]]  
Out[123]:  
b      1.0  
d      3.0  
dtype: float64
```

Selection with loc and iloc

```
In [137]: data.loc['Colorado', ['two', 'three']]
```

```
Out[137]:
```

```
two      5
```

```
three    6
```

```
Name: Colorado, dtype: int64
```

```
In [138]: data.iloc[2, [3, 0, 1]]
```

```
Out[138]:
```

```
four     11
```

```
one      8
```

```
two      9
```

```
Name: Utah, dtype: int64
```

```
In [139]: data.iloc[2]
```

```
Out[139]:
```

```
one      8
```

```
two      9
```

```
three   10
```

Apply function

```
In [193]: f = lambda x: x.max() - x.min()
```

```
In [194]: frame.apply(f)
```

```
Out[194]:
```

```
b    1.802165
```

```
d    1.684034
```

```
e    2.689627
```

```
dtype: float64
```

Function application and mapping

```
In [192]: np.abs(frame)
```

```
Out[192]:
```

	b	d	e
Utah	0.204708	0.478943	0.519439
Ohio	0.555730	1.965781	1.393406
Texas	0.092908	0.281746	0.769023
Oregon	1.246435	1.007189	1.296221

Descriptive statistics

```
In [233]: df.sum(axis='columns')
```

```
Out[233]:
```

```
a    1.40
```

```
b    2.60
```

```
c    NaN
```

```
d   -0.55
```

```
dtype: float64
```

```
In [235]: df.idxmax()
```

```
Out[235]:
```

```
one    b
```

```
two    d
```

```
dtype: object
```

```
In [234]: df.mean(axis='columns', skipna=False)
```

```
Out[234]:
```

```
a    NaN
```

```
b    1.300
```

```
c    NaN
```

```
In [237]: df.describe()
```

```
Out[237]:
```

	one	two
count	3.000000	2.000000
mean	3.083333	-2.900000
std	3.493685	2.262742
min	0.750000	-4.500000
25%	1.075000	-3.700000
50%	1.400000	-2.900000
75%	4.250000	-2.100000
max	7.100000	-1.300000

```
In [236]: df.cumsum()
```

```
Out[236]:
```

	one	two
a	1.40	NaN
b	8.50	-4.5
c	NaN	NaN
d	9.25	-5.8

Method	Description
count	Number of non-NA values
describe	Compute set of summary statistics for Series or each DataFrame column
min, max	Compute minimum and maximum values
argmin, argmax	Compute index locations (integers) at which minimum or maximum value obtained, respectively
idxmin, idxmax	Compute index labels at which minimum or maximum value obtained, respectively
quantile	Compute sample quantile ranging from 0 to 1
sum	Sum of values
mean	Mean of values
median	Arithmetic median (50% quantile) of values
mad	Mean absolute deviation from mean value
prod	Product of all values
var	Sample variance of values
std	Sample standard deviation of values
skew	Sample skewness (third moment) of values
kurt	Sample kurtosis (fourth moment) of values
cumsum	Cumulative sum of values
cummin, cummax	Cumulative minimum or maximum of values, respectively
cumprod	Cumulative product of values
diff	Compute first arithmetic difference (useful for time series)
pct_change	Compute percent changes

Reading and writing data

```
In [9]: df = pd.read_csv('examples/ex1.csv')
```

```
In [10]: df
```

```
Out[10]:
```

	a	b	c	d	message
0	1	2	3	4	hello
1	5	6	7	8	world
2	9	10	11	12	foo

```
In [13]: pd.read_csv('examples/ex2.csv', header=None)
```

```
Out[13]:
```

	0	1	2	3	4
0	1	2	3	4	hello
1	5	6	7	8	world
2	9	10	11	12	foo

Filtering out missing values

```
In [19]: data = pd.DataFrame([[1., 6.5, 3.], [1., NA, NA],  
.....:                      [NA, NA, NA], [NA, 6.5, 3.]])
```

```
In [20]: cleaned = data.dropna()
```

```
In [21]: data
```

```
Out[21]:
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
2	NaN	NaN	NaN
3	NaN	6.5	3.0

```
In [22]: cleaned
```

```
Out[22]:
```

	0	1	2
0	1.0	6.5	3.0

Passing `how='all'` will only drop rows that are all NA:

```
In [23]: data.dropna(how='all')
```

```
Out[23]:
```

```
    0    1    2
```

Filling missing data

```
In [33]: df.fillna(0)
```

```
Out[33]:
```

Reading data in pandas

Table 6-1. Parsing functions in pandas

Function	Description
<code>read_csv</code>	Load delimited data from a file, URL, or file-like object; use comma as default delimiter
<code>read_table</code>	Load delimited data from a file, URL, or file-like object; use tab (<code>'\t'</code>) as default delimiter
<code>read_fwf</code>	Read data in fixed-width column format (i.e., no delimiters)
<code>read_clipboard</code>	Version of <code>read_table</code> that reads data from the clipboard; useful for converting tables from web pages
<code>read_excel</code>	Read tabular data from an Excel XLS or XLSX file
<code>read_hdf</code>	Read HDF5 files written by pandas
<code>read_html</code>	Read all tables found in the given HTML document
<code>read_json</code>	Read data from a JSON (JavaScript Object Notation) string representation
<code>read_msgpack</code>	Read pandas data encoded using the MessagePack binary format
<code>read_pickle</code>	Read an arbitrary object stored in Python pickle format
<code>read_sas</code>	Read a SAS dataset stored in one of the SAS system's custom storage formats
<code>read_sql</code>	Read the results of a SQL query (using SQLAlchemy) as a pandas DataFrame
<code>read_stata</code>	Read a dataset from Stata file format
<code>read_feather</code>	Read the Feather binary file format

Optional arguments

- Indexing Can treat one or more columns as the returned DataFrame, and whether to get column names from the file, the user, or not at all. Type inference and data conversion This includes the user-defined value conversions and custom list of missing value markers.
- Datetime parsing Includes combining capability, including combining date and time information spread over multiple columns into a single column in the result.
- Iterating Support for iterating over chunks of very large files.
- Unclean data issues Skipping rows or a footer, comments, or other minor things like numeric data with thousands separated by commas.

Examples

```
df = pd.read_csv('examples/ex1.csv')
```

```
pd.read_table('examples/ex1.csv', sep=',')
```

```
pd.read_csv('examples/ex2.csv', header=None)
```

```
pd.read_csv('examples/ex2.csv', names=['a', 'b', 'c', 'd', 'message'])
```

```
names = ['a', 'b', 'c', 'd', 'message']
```

```
pd.read_csv('examples/ex2.csv', names=names, index_col='message')
```


- Passing regular expression as separator
 - `result = pd.read_table('examples/ex3.txt', sep='\s+')`
- Skipping rows
 - `pd.read_csv('examples/ex4.csv', skiprows=[0, 2, 3])`

Handling null values

```
result = pd.read_csv('examples/ex5.csv', na_values=['NULL'])
```

```
In [31]: sentinels = {'message': ['foo', 'NA'], 'something': ['two']}
```

```
In [32]: pd.read_csv('examples/ex5.csv', na_values=sentinels)
```

Table 6-2. Some read_csv/read_table function arguments

Argument	Description
path	String indicating filesystem location, URL, or file-like object
sep or delimiter	Character sequence or regular expression to use to split fields in each row
header	Row number to use as column names; defaults to 0 (first row), but should be None if there is no header row
index_col	Column numbers or names to use as the row index in the result; can be a single name/number or a list of them for a hierarchical index
names	List of column names for result, combine with header=None

Argument	Description
<code>skiprows</code>	Number of rows at beginning of file to ignore or list of row numbers (starting from 0) to skip.
<code>na_values</code>	Sequence of values to replace with NA.
<code>comment</code>	Character(s) to split comments off the end of lines.
<code>parse_dates</code>	Attempt to parse data to <code>datetime</code> ; <code>False</code> by default. If <code>True</code> , will attempt to parse all columns. Otherwise can specify a list of column numbers or name to parse. If element of list is tuple or list, will combine multiple columns together and parse to date (e.g., if date/time split across two columns).
<code>keep_date_col</code>	If joining columns to parse date, keep the joined columns; <code>False</code> by default.
<code>converters</code>	Dict containing column number or name mapping to functions (e.g., <code>{ 'foo' : f }</code> would apply the function <code>f</code> to all values in the <code>'foo'</code> column).
<code>dayfirst</code>	When parsing potentially ambiguous dates, treat as international format (e.g., <code>7/6/2012</code> -> June 7, 2012); <code>False</code> by default.
<code>date_parser</code>	Function to use to parse dates.
<code>nrows</code>	Number of rows to read from beginning of file.
<code>iterator</code>	Return a <code>TextParser</code> object for reading file piecemeal.
<code>chunksize</code>	For iteration, size of file chunks.
<code>skip_footer</code>	Number of lines to ignore at end of file.
<code>verbose</code>	Print various parser output information, like the number of missing values placed in non-numeric columns.
<code>encoding</code>	Text encoding for Unicode (e.g., <code>'utf-8'</code> for UTF-8 encoded text).
<code>squeeze</code>	If the parsed data only contains one column, return a Series.
<code>thousands</code>	Separator for thousands (e.g., <code>' , '</code> or <code>' . '</code>).

Reading text file in pieces

```
In [33]: pd.options.display.max_rows = 10
```

```
In [36]: pd.read_csv('examples/ex6.csv', nrows=5)
```

```
chunker = pd.read_csv('examples/ex6.csv', chunksize=1000)
```

Iterating over chunk

```
chunker = pd.read_csv('examples/ex6.csv', chunksize=1000)

tot = pd.Series([])
for piece in chunker:
    tot = tot.add(piece['key'].value_counts(), fill_value=0)

tot = tot.sort_values(ascending=False)
```

Writing file

```
data.to_csv('examples/out.csv')
```

```
import sys
```

```
data.to_csv(sys.stdout, sep='|')
```

```
data.to_csv(sys.stdout, na_rep='NULL')
```

```
data.to_csv(sys.stdout, index=False, header=False)
```

```
data.to_csv(sys.stdout, index=False, columns=['a', 'b', 'c'])
```


Working with delimited format

```
import csv
f = open('examples/ex7.csv')

reader = csv.reader(f)

with open('examples/ex7.csv') as f:
    lines = list(csv.reader(f))

header, values = lines[0], lines[1:]
```

Writing csv data

```
with open('mydata.csv', 'w') as f:  
    writer = csv.writer(f, dialect=my_dialect)  
    writer.writerow(('one', 'two', 'three'))  
    writer.writerow(('1', '2', '3'))  
    writer.writerow(('4', '5', '6'))  
    writer.writerow(('7', '8', '9'))
```

JSON format

```
obj = """
{"name": "Wes",
 "places_lived": ["United States", "Spain", "Germany"],
 "pet": null,
 "siblings": [{"name": "Scott", "age": 30, "pets": ["Zeus", "Zuko"]},
               {"name": "Katie", "age": 38,
                "pets": ["Sixes", "Stache", "Cisco"]}]}
"""
```

- The `pandas.read_json` can automatically convert JSON datasets in specific arrangements into a Series or DataFrame
- The default options for `pandas.read_json` assume that each object in the JSON array is a row in the table

```
data = pd.read_json('examples/example.json')
```

```
print(data.to_json())
```

```
print(data.to_json(orient='records'))
```

Web scraping

```
conda install lxml  
pip install beautifulsoup4 html5lib
```

- The `pandas.read_html` function has a number of options, but by default it searches for and attempts to parse all tabular data contained within tags.

```
In [73]: tables = pd.read_html('examples/fdic_failed_bank_list.html')
```

```
In [74]: len(tables)
```

```
Out[74]: 1
```

```
In [75]: failures = tables[0]
```

```
In [76]: failures.head()
```

```
Out[76]:
```

	Bank Name	City	ST	CERT	\
0	Allied Bank	Mulberry	AR	91	
1	The Woodbury Banking Company	Woodbury	GA	11297	
2	First CornerStone Bank	King of Prussia	PA	35312	

Saving in binary format

```
frame.to_pickle('examples/frame_pickle')
```

```
pd.read_pickle('examples/frame_pickle')
```

HD5 format

```
In [92]: frame = pd.DataFrame({'a': np.random.randn(100)})
```

```
In [93]: store = pd.HDFStore('mydata.h5')
```

```
In [94]: store['obj1'] = frame
```

```
In [95]: store['obj1_col'] = frame['a']
```

```
In [96]: store
```


Excel format

```
xlsx = pd.ExcelFile('examples/ex1.xlsx')
```

```
pd.read_excel(xlsx, 'Sheet1')
```

```
writer = pd.ExcelWriter('examples/ex2.xlsx')
```

```
frame.to_excel(writer, 'Sheet1')
```

```
writer.save()
```

Interacting with Web API

```
import requests
```

```
url = 'https://api.github.com/repos/pandas-dev/pandas/issues'
```

```
resp = requests.get(url)
```

```
resp
```

```
<Response [200]>
```

```
data = resp.json()
```

```
data[0]['title']
```


Interacting with database

```
In [121]: import sqlite3
```

```
In [122]: query = """
.....: CREATE TABLE test
.....: (a VARCHAR(20), b VARCHAR(20),
.....:  c REAL,          d INTEGER
.....: );"""
```

```
In [123]: con = sqlite3.connect('mydata.sqlite')
```

```
In [124]: con.execute(query)
```

```
Out[124]: <sqlite3.Cursor at 0x7f6b12a50f10>
```

```
In [125]: con.commit()
```

```
data = [('Atlanta', 'Georgia', 1.25, 6),  
        ('Tallahassee', 'Florida', 2.6, 3),  
        ('Sacramento', 'California', 1.7, 5)]
```

```
stmt = "INSERT INTO test VALUES(?, ?, ?, ?)"
```

```
con.executemany(stmt, data)  
<sqlite3.Cursor at 0x7f6b15c66ce0>
```

```
con.commit()
```

```
cursor = con.execute('select * from test')
```

```
rows = cursor.fetchall()
```

```
In [133]: cursor.description
```

```
Out[133]:
```

```
(( 'a', None, None, None, None, None, None),  
 ( 'b', None, None, None, None, None, None),  
 ( 'c', None, None, None, None, None, None),  
 ( 'd', None, None, None, None, None, None))
```

```
In [134]: pd.DataFrame(rows, columns=[x[0] for x in cursor.description])
```

```
Out[134]:
```

Data Cleaning and Preparation

Handling missing data

```
In [10]: string_data = pd.Series(['aardvark', 'artichoke', np.nan, 'avocado'])
```

```
In [11]: string_data
```

```
Out[11]:
```

```
0    aardvark
1    artichoke
2         NaN
3     avocado
dtype: object
```

```
In [12]: string_data.isnull()
```

```
Out[12]:
```

```
0    False
1    False
2     True
3    False
dtype: bool
```


Table 7-1. NA handling methods

Argument	Description
<code>dropna</code>	Filter axis labels based on whether values for each label have missing data, with varying thresholds for how much missing data to tolerate.
<code>fillna</code>	Fill in missing data with some value or using an interpolation method such as <code>'ffill'</code> or <code>'bfill'</code> .
<code>isnull</code>	Return boolean values indicating which values are missing/NA.
<code>notnull</code>	Negation of <code>isnull</code> .

```
In [15]: from numpy import nan as NA
```

```
In [16]: data = pd.Series([1, NA, 3.5, NA, 7])
```

```
In [17]: data.dropna()
```

```
Out[17]:
```

```
0    1.0
```

```
2    3.5
```

```
4    7.0
```

```
dtype: float64
```

```
In [18]: data[data.notnull()]
```

```
Out[18]:
```

```
0    1.0
```

```
2    3.5
```

```
4    7.0
```

```
dtype: float64
```

dropna options

- With DataFrame objects, things are a bit more complex. You may want to drop rows or columns that are all NA or only those containing any NAs. dropna by default drops any row containing a missing value
- Passing how='all' will only drop rows that are all NA
- To drop columns in the same way, pass axis=1
- df.dropna(thresh=2)

Filling In Missing Data

- Calling `fillna` with a constant replaces missing values with that value
- Calling `fillna` with a dict, you can use a different fill value for each column
- `df.fillna({1: 0.5, 2: 0})`
- `fillna` returns a new object, but you can modify the existing object in-place:
 - `_ = df.fillna(0, inplace=True)`

```
In [37]: df = pd.DataFrame(np.random.randn(6, 3))
```

```
In [38]: df.iloc[2:, 1] = NA
```

```
In [39]: df.iloc[4:, 2] = NA
```

```
In [40]: df
```

```
Out[40]:
```

	0	1	2
0	0.476985	3.248944	-1.021228
1	-0.577087	0.124121	0.302614
2	0.523772	NaN	1.343810
3	-0.713544	NaN	-2.370232
4	-1.860761	NaN	NaN
5	-1.265934	NaN	NaN

```
In [41]: df.fillna(method='ffill')
```

```
Out[41]:
```

	0	1	2
0	0.476985	3.248944	-1.021228
1	-0.577087	0.124121	0.302614
2	0.523772	0.124121	1.343810
3	-0.713544	0.124121	-2.370232
4	-1.860761	0.124121	-2.370232
5	-1.265934	0.124121	-2.370232

```
In [42]: df.fillna(method='ffill', limit=2)
```

```
Out[42]:
```

	0	1	2
0	0.476985	3.248944	-1.021228
1	-0.577087	0.124121	0.302614
2	0.523772	0.124121	1.343810
3	-0.713544	0.124121	-2.370232
4	-1.860761	NaN	-2.370232
5	-1.265934	NaN	-2.370232

Table 7-2. fillna function arguments

Argument	Description
value	Scalar value or dict-like object to use to fill missing values
method	Interpolation; by default 'ffill' if function called with no other arguments
axis	Axis to fill on; default axis=0
inplace	Modify the calling object without producing a copy
limit	For forward and backward filling, maximum number of consecutive periods to fill

Removing Duplicates

- `data.duplicated()`
- Relatedly, `drop_duplicates` returns a DataFrame where the duplicated array is False
- You can specify any subset of them to detect duplicates
- `data.drop_duplicates(['k1'])`
- `duplicated` and `drop_duplicates` by default keep the first observed value combination. Passing `keep='last'` will return the last one

Map

```
In [52]: data = pd.DataFrame({'food': ['bacon', 'pulled pork', 'bacon',  
.....:                               'Pastrami', 'corned beef', 'Bacon',  
.....:                               'pastrami', 'honey ham', 'nova lox'],  
.....:                       'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})
```

```
lowercased = data['food'].str.lower()
```

```
In [57]: data['animal'] = lowercased.map(meat_to_animal)
```

```
In [59]: data['food'].map(lambda x: meat_to_animal[x.lower()])
```

```
meat_to_animal = {  
    'bacon': 'pig',  
    'pulled pork': 'pig',  
    'pastrami': 'cow',  
    'corned beef': 'cow',  
    'honey ham': 'pig',  
    'nova lox': 'salmon'  
}
```


Replace

- `data.replace(-999, np.nan)`
- `data.replace([-999, -1000], np.nan)`
- `data.replace([-999, -1000], [np.nan, 0])`
- `data.replace({-999: np.nan, -1000: 0})`

Map with index

```
In [66]: data = pd.DataFrame(np.arange(12).reshape((3, 4)),  
.....:                      index=['Ohio', 'Colorado', 'New York'],  
.....:                      columns=['one', 'two', 'three', 'four'])
```

```
In [67]: transform = lambda x: x[:4].upper()
```

```
In [68]: data.index.map(transform)
```

```
Out[68]: Index(['OHIO', 'COLO', 'NEW '], dtype='object')
```

```
In [69]: data.index = data.index.map(transform)
```

Other index methods

- `data.rename(index=str.title, columns=str.upper)`
- `data.rename(index={'OHIO': 'INDIANA'}, columns={'three': 'peekaboo'})`
- `data.rename(index={'OHIO': 'INDIANA'}, inplace=True)`

Discretization and binning

- Let's divide data into bins
 - `ages = [20, 22, 25, 27, 21, 23, 37, 31, 61, 45, 41, 32]`
 - `bins = [18, 25, 35, 60, 100]`
 - `cats = pd.cut(ages, bins)`
- Lets see the code for each data and categories:
 - `cats.codes`
 - `cats.categories`
- What are the counts for each category?
 - `pd.value_counts(cats)`

- Making the left side a close interval
 - `pd.cut(ages, [18, 26, 36, 61, 100], right=False)`
- Providing your own group names:
 - `group_names = ['Youth', 'YoungAdult', 'MiddleAged', 'Senior']`
 - `pd.cut(ages, bins, labels=group_names)`

Detecting and filtering outliers

- `data.describe()`
- Suppose you wanted to find values in one of the columns exceeding 3 in absolute value:
 - `col = data[2]`
 - `col[np.abs(col) > 3]`
- To select all rows having a value exceeding 3 or -3 , you can use the `any` method on a boolean DataFrame:
 - `data[(np.abs(data) > 3).any(1)]`
- Here is code to cap values outside the interval -3 to 3
 - `data[np.abs(data) > 3] = np.sign(data) * 3`

Permutation and random sampling

- Calling permutation with the length of the axis you want to permute produces an array of integers indicating the new ordering:
- `df = pd.DataFrame(np.arange(5 * 4).reshape((5, 4)))`
- `sampler = np.random.permutation(5)`
- `df.take(sampler)`
- To select a random subset without replacement, you can use the sample method on Series and DataFrame:
 - `df.sample(n=3)`

- To generate a sample with replacement (to allow repeat choices), pass `replace=True` to `sample`:
 - `choices = pd.Series([5, 7, -1, 6, 4])`
 - `draws = choices.sample(n=10, replace=True)`